



CONTENTS

1	Projections	3
1.1	Projections in Python	6
1.2	Where to Learn More	6
2	The Gram-Schmidt Process	7
2.1	The Process	7
2.2	Example Calculation	8
2.3	The Gram-Schmidt Process in Python	10
2.4	Where to Learn More	11
3	Eigenvectors and Eigenvalues	13
3.1	Definition	14
3.2	Finding Eigenvalues and Eigenvectors	14
3.3	Example	14
3.4	Eigenvalues and Eigenvectors in Python	17
3.5	Where to Learn More	17
4	Singular Value Decomposition	19
4.1	Definition	19
4.2	Applications of SVD	19
4.3	Calculating SVD Manually	20
4.4	Singular Value Decomposition with Python	23
4.5	Sign Ambiguity	24
4.6	SVD Applied to Image Compression	24

4.7	Where to Learn More	25
5	Tackling Difficult Problems: Positive Semidefinite Matrices	27
5.1	NP Problems	27
5.2	A Past Approach: Minimizing Errors in Neural Networks	28
5.3	Positive Definite and Semidefinite Matrices	29
5.4	Identifying and Constructing a Positive Semidefinite Matrix	30
5.5	The Max Cut Problem	31
5.6	The Max Cut Problem Solved in Python	33
A	Answers to Exercises	35
	Index	37

Projections


The word projection has two main meanings in everyday life. One is a projection as a forecast or estimate of something in the future based on the current situation. Another is the result of shining a light to cast a shadow or show a movie. Both these definitions apply to mathematical projection.

Projections are used in many fields such as science, math, engineering, and finance. Here are a few examples:


- Investors evaluate risk and return of a portfolio by projecting an asset's return onto a reference portfolio.
- Astronomers analyze the motion of stellar objects by projecting the object's true motion onto the plane of the sky.
- Robotics engineers use projections to prevent robots from running into obstacles by projecting the robot's position onto the optimal path.

Mathematically, a projection describes the relationship of one vector to another in terms of direction and orthogonality. Given two vectors, \mathbf{u} and \mathbf{v} , the projection of \mathbf{u} onto \mathbf{v} separates \mathbf{u} into two components. The first component signifies how much \mathbf{u} lies in the direction of \mathbf{v} . The second signifies the component of \mathbf{u} that is orthogonal (perpendicular) to \mathbf{v} .


The figure depicts a projection. The perpendicular line dropped from the end of \mathbf{u} is the orthogonal component. The portion of \mathbf{u} that lies in the direction of \mathbf{v} is the blue segment.



You can also think of a projection as the shadow cast by one vector onto each other by an overhead light.



The projected vector can be in any direction and its length can extend beyond the vector onto which it is projecting.



To calculate the projection of \mathbf{v} onto \mathbf{u} , use this formula:

$$\text{proj}_{\mathbf{v}}(\mathbf{u}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{v}\|^2} \mathbf{v}$$

Note that the denominator is the magnitude squared of vector \mathbf{v} .

$$(\sqrt{a_1^2 + a_2^2 + \dots + a_n^2})^2$$

You learned previously that this is the same as the dot product of a vector with itself.

$$\mathbf{v} \cdot \mathbf{v}$$

In the examples that follow, we'll simplify to the dot product notation.

Let's look at a specific example:

$$\mathbf{u} = (1, 4, 6)$$

$$\mathbf{v} = (-2, 6, 2)$$

$$\text{proj}_{\mathbf{v}}(\mathbf{u}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{v}\|^2} \mathbf{v}$$

$$\text{proj}_{\mathbf{v}}(\mathbf{u}) = \frac{(1, 4, 6) \cdot (-2, 6, 2)}{(-2, 6, 2) \cdot (-2, 6, 2)} (-2, 6, 2)$$

$$\text{proj}_{\mathbf{v}}(\mathbf{u}) = \left(\frac{34}{44}\right) (-2, 6, 2)$$

$$\text{proj}_{\mathbf{v}}(\mathbf{u}) = (-1.545, 4.64, 1.545)$$

As you work your way through this course, you'll have a chance to apply the calculations you learn in this chapter to a variety of problems. Specifically, the next chapter shows how to transform a set of linearly independent vectors into a set of orthogonal ones. Projections are essential to that transformation.

Exercise 1 Projections

Find the projection of **a** on **b** where:

$$\mathbf{a} = (1, 3)$$

$$\mathbf{b} = (-4, 6)$$

Working Space

Answer on Page 35

1.1 Projections in Python

Create a file called `vectors_projections.py` and enter this code:

```
import numpy as np

# define two vectors
a = np.array([1, 4, 6])
b = np.array([-2, 6, 2])

# use np.dot() to calculate the dot product
projection_a_on_b = (np.dot(a, b)/np.dot(b, b))*b

print("The projection of vector a on vector b is:", projection_a_on_b)
```

1.2 Where to Learn More

Watch this Introduction to Projections from Khan Academy <https://rb.gy/yf0i3>

The Gram-Schmidt Process

The Gram-Schmidt process is a method use to transform a set of linearly independent vectors to a set of orthogonal (perpendicular vectors). The original vectors and the transformed vectors span the same subspace.

The process was named after two mathematicians: Jørgen Pedersen Gram, a Danish actuary mathematician, and Erhard Schmidt, a German mathematician. The men developed the orthogonalization process independently. Gram introduced the process in 1883 whereas Schmidt did his work in 1907. It wasn't named the Gram-Schmidt process until sometime later, after both mathematicians became well-known in the mathematical community.

In the last chapter you learned how to calculate a projection of one vector on another. Given two vectors, u and v , the projection separates u into the part that is orthogonal to v and the part that has a dependency with v . The Gram-Schmidt process builds on the notion of projections to iteratively strip away dependencies between vectors until the vectors that remain are orthogonal. If there happens to be a vector that is a linear combination of the others, that vector will be reduced to the zero vector. The vectors that remain define a new basis for space spanned by the original set of vectors.

Gram-Schmidt has many practical applications in science and engineering. These are two examples:

1. In signal processing, it can represent an audio signal with fewer components making it easier to isolate and remove noise.
2. In statistics and data analysis, it can reduce the complexity of a dataset so that it is easier to see which aspects or features contribute to the analysis.

2.1 The Process

The Gram-Schmidt process orthonormalizes a set of vectors in an inner product space, most commonly the Euclidean space \mathbb{R}^n . The process takes a finite, linearly independent set $S = \{v_1, v_2, \dots, v_k\}$ for $k \leq n$, and generates an orthogonal set $S' = \{u_1, u_2, \dots, u_k\}$ that spans the same k -dimensional subspace of \mathbb{R}^n as S .

Let's look at how the process works. Given a set of vectors $S = \{v_1, v_2, \dots, v_k\}$, the Gram-Schmidt process is as follows:

1. Let $u_1 = v_1$.
2. For $j = 2, 3, \dots, k$:
 - (a) Let $w_j = v_j - \sum_{i=1}^{j-1} \frac{\langle v_j, u_i \rangle}{\langle u_i, u_i \rangle} u_i$
 - (b) Let $u_j = w_j$

Here, $\langle \cdot, \cdot \rangle$ denotes the inner product.

The set of vectors $S' = \{u_1, u_2, \dots, u_k\}$ obtained from this process is orthogonal, but not necessarily orthonormal. To create an orthonormal set, you simply need to normalize each vector u_i to unit length. That is, $u'_i = \frac{u_i}{\|u_i\|}$, where $\|\cdot\|$ denotes the norm (or length) of a vector.

Among other things, making vectors orthonormal simplifies calculations, makes it easier to define rotations and transformations, and provides a framework for calculations in fields such as quantum mechanics.

2.2 Example Calculation

Given a set of linearly independent vectors, we will use the Gram-Schmidt process to find an orthogonal basis.

Let

$$W = \text{Span}(x_1, x_2, x_3)$$

where

$$x_1 = (1, 2, -2)$$

$$x_2 = (1, 0, -4)$$

$$x_3 = (5, 2, 0)$$

The three orthogonal vectors will define the same subspace as the original vectors.

The first vector of the orthogonal subspace is easy to define. We set it to be the same as x_1 .

$$v_1 = x_1 = (1, 2, -2)$$

The second orthogonal vector is a projection of x_2 onto v_1 . You learned projections in the last chapter, so this should be fairly straightforward.

$$v_2 = x_2 - \frac{x_2 v_1}{v_1 v_1} v_1$$

Substitute the values:

$$v_2 = (1, 0, -4) - \frac{(1, 0, -4)(1, 2, -2)}{(1, 2, -2)(1, 2, -2)} (1, 2, -2)$$

Calculate the coefficient for v_1 :

$$v_2 = (1, 0, -4) - \frac{9}{9} (1, 2, -2)$$

Perform the subtraction:

$$v_2 = (0, -2, -2)$$

The third vector for the orthogonal subspace is a projection onto v_1 and v_2 .

$$v_3 = x_3 - \frac{x_3 v_1}{v_1 v_1} v_1 - \frac{x_3 v_2}{v_2 v_2} v_2$$

Substitute the values:

$$v_3 = (5, 2, 0) - \frac{(5, 2, 0)(1, 2, -2)}{(1, 2, -2)(1, 2, -2)} (1, 2, -2) - \frac{(5, 2, 0)(1, 0, -4)}{(1, 0, -4)(1, 0, -4)} (1, 0, -4)$$

$$v_3 = (5, 2, 0) - (9/9)(1, 2, -2) - (-4/8)(1, 0, -4)$$

$$v_3 = (5, 2, 0) - (1, 2, -2) + (1/2)(1, 0, -4)$$

$$v_3 = (5, 2, 0) - (1, 2, -2) + (0, -1, -1)$$

$$v_3 = (4, -1, 1)$$

This set of vectors is orthogonal, so we need to normalize them so that the vectors are orthonormal. Recall that an orthonormal vector has a length of 1 and is computed using this formula:

$$\text{normalizedVector} = \text{vector} / \text{np.sqrt}(\text{np.sum}(\text{vector} * * 2))$$

Thus the normalized set of vectors is:

$$v_1 = (0.33, 0.67, -0.67)$$

$$v_2 = (0.0, -0.71, -0.71)$$

$$v_3 = (0.94, -0.24, 0.24)$$

Exercise 2 Gram-Schmidt Process

Use the Gram-Schmidt process to find an orthogonal basis for the span defined by x_1, x_2 where:

$$x_1 = (1, 1, 1)$$

$$x_2 = (0, 1, 1)$$

Working Space

Answer on Page 35

2.3 The Gram-Schmidt Process in Python

Create a file called `vectors_gram-schmidt.py` and enter this code:

```
# import numpy to perform operations on vector
import numpy as np

# Find an orthogonal basis for the span of these three vectors
x1 = np.array([1, 2, -2])
x2 = np.array([1, 0, -4])
x3 = np.array([5, 2, 0])

# v1 = x1
v1 = x1
print("v1 = ", v1)

# v2 = x2 - (the projection of x2 on v1)
v2 = x2 - (np.dot(x2, v1) / np.dot(v1, v1)) * v1
print("v2 = ", v2)

# v3 = x3 - (the projection of x3 on v1) - (the projection of x3 on v2)
v3 = x3 - (np.dot(x3, v1) / np.dot(v1, v1)) * v1 - (np.dot(x3, v2) / np.dot(v2, v2)) * v2
print("v3 = ", v3)

# Next, normalize each vector to get a set of vectors that is both orthogonal and orthonormal:
v1_norm = v1 / np.sqrt(np.sum(v1**2))
v2_norm = v2 / np.sqrt(np.sum(v2**2))
v3_norm = v3 / np.sqrt(np.sum(v3**2))
```

```
print("v1_norm = ", v1_norm)
print("v2_norm = ", v2_norm)
print("v3_norm = ", v3_norm)
```

2.4 Where to Learn More


Watch this video from Khan Academy about the Gram-Schmidt process: <https://www.khanacademy.org/math/linear-algebra/alternate-bases/orthonormal-basis/v/linear-algebra-t>

Eigenvectors and Eigenvalues

Like many specialized disciplines, Linear Algebra uses many unfamiliar terms whose origin you might wonder about. Eigenvectors and eigenvalues are two of them. If you know German, you'll recognize that *eigen* means inherent or a characteristic attribute. Named by the German mathematician David Hilbert, an eigenvector mathematically describes a characteristic feature of an object, that remains unchanged after transformation. You can think of an eigenvector as the direction that doesn't change direction. An eigenvector characterizes a linear transformation whereas its eigenvalue tells how much the vector is scaled. Eigenvalues can be negative or positive. A negative value indicates the direction of the eigenvector is reversed.


Eigenvalues and eigenvectors are a way to break down matrices that can simplify many calculations and enable us to understand various properties of the matrix. They are widely used in physics and engineering for stability analysis, vibration analysis, and many other applications.

Let's look at a visual example.



You can see that the image on the right is a skewed version of the image on the left. Look closely at the vectors and you'll notice that the one of the vectors is pointing in the same direction in both images, while the direction of the other two vectors has changed. The eigenvector is the one at the bottom that points 0 degrees (or you can think of due east) in both images. Thus the characteristic attribute of both images is their horizontal direction.

When you overlay the vectors from one image over the other, you'll notice that the horizontal vector, while the same direction in both images, is a bit longer in the skewed version. The scale of the stretch is described by an eigenvalue.



3.1 Definition

Given a square matrix A , a non-zero vector v is an eigenvector of A if multiplying A by v results in a scalar multiple of v , i.e.,

$$Av = \lambda v \quad (3.1)$$

where λ is a scalar known as the eigenvalue corresponding to the eigenvector v .

3.2 Finding Eigenvalues and Eigenvectors

You find the eigenvalues of a matrix A by solving the characteristic equation:

$$\det(A - \lambda I) = 0 \quad (3.2)$$

where $\det(\cdot)$ denotes the determinant, I is the identity matrix of the same size as A , and λ is a scalar.

Once you find the eigenvalues, you can find the corresponding eigenvectors by substituting each eigenvalue into the equation $Av = \lambda v$, and solving for v .

3.3 Example

For a 2×2 matrix $A = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$, the characteristic equation is:

$$(a - \lambda)(d - \lambda) - bc = 0 \quad (3.3)$$

Solving this equation gives the eigenvalues. Substituting each eigenvalue back into the equation $Av = \lambda v$ gives the corresponding eigenvectors.

Let matrix $A =$

$$\begin{bmatrix} 5 & 4 \\ 1 & 2 \end{bmatrix}$$

The characteristic equation is:

$$|A - \lambda I| = 0$$

$$\begin{bmatrix} 5 - \lambda & 4 \\ 1 & 2 - \lambda \end{bmatrix} = 0$$

$$(5 - \lambda)(2 - \lambda) - (4)(1) = 0$$

$$10 - 5\lambda - 2\lambda + \lambda^2 - 4 = 0$$

$$\lambda^2 - 7\lambda + 6 = 0$$

$$(\lambda - 6)(\lambda - 1) = 0$$

$$\lambda = 6, \lambda = 1$$

Now that you have the eigen values you can substitute these values into the equation:

$$|A - \lambda I| = 0$$

For $\lambda = 1$:

$$(A - \lambda I)v = 0$$

$$\begin{bmatrix} 5 - 1 & 4 \\ 1 & 2 - 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} 4 & 4 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Next, use elementary row transformation by multiplying row 2 by 4 and then subtracting row 1.

$$\begin{bmatrix} 4 & 4 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Now you can expand as an equation:

$$4x + 4y = 0$$

Assume $y = w$

$$4x = -4w$$

$$x = -w$$

The solution is:

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -w \\ w \end{bmatrix} = w \begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

So the eigenvector is:

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

Now we need to substitute the other eigenvalue, 6, into the equation and follow the same procedure for finding the eigenvector.

$$\begin{bmatrix} 5-6 & 4 \\ 1 & 2-6 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} -1 & 4 \\ 1 & -4 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Next, use elementary row transformation by adding row 1 to row 2.

$$\begin{bmatrix} -1 & 4 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Expand as an equation:

$$-x + 4y = 0$$

$$\text{Assume } y = w$$

$$-x + 4w = 0$$

$$x = 4w$$

The solution is:

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 4w \\ w \end{bmatrix} = w \begin{bmatrix} 4 \\ 1 \end{bmatrix}$$

So the eigenvector is:

$$\begin{bmatrix} 4 \\ 1 \end{bmatrix}$$

In conclusion, the eigenvectors of the given 2 x 2 matrix are:

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix} \text{ and } \begin{bmatrix} 4 \\ 1 \end{bmatrix}$$

3.4 Eigenvalues and Eigenvectors in Python

Create a file called `vectors_eigen.py` and enter this code:

```
# import numpy to perform operations on vector
import numpy as np
from numpy.linalg import eig

a = np.array([[2, 2, 4],
              [1, 3, 5],
              [2, 3, 4]])
eigenvalue, eigenvector = eig(a)

# The values are not in any particular order
print('Eigenvalues:', eigenvalue)

# The eig function returns the normalize vectors
print('Eigenvectors:', eigenvector)
```

3.5 Where to Learn More

Watch this video from Khan Academy, *Introduction to Eigenvectors*: <https://rb.gy/mse7i>

Singular Value Decomposition

In the previous chapter you learned how to calculate eigenvalues and eigenvectors. But not every matrix has them. For those matrices, singular values and singular vectors are analogous features.

Singular Value Decomposition (SVD) is a matrix factorization technique that breaks down a matrix into three matrices that represent the structure and properties of the original matrix. The decomposed matrices make calculations easier and provide insight into the original matrix. Basically, SVD can transform a high dimension, highly variable set of data into a set of uncorrelated data points that reveal subgroupings that you might not have noticed in the original data. SVD tells us that a linear transformation can be thought of as a rotation, scaling, and another rotation.

4.1 Definition

For any $m \times n$ matrix A , SVD decomposes the matrix into three matrices.

$$A = U\Sigma V^T \quad (4.1)$$

- U is an orthogonal matrix whose size is $m \times m$. Its columns are the eigenvectors of AA^T . These are the left singular vectors of A . Because U is orthogonal, $U^T U = I$
- V is an orthogonal matrix whose size is $n \times n$ matrix. Its columns are the eigenvectors of $A^T A$. These are the right singular vectors of A . Because V is orthogonal, $V^T V = I$.
- Σ is a diagonal matrix that is the same size as A . Its diagonal contains the singular values of A , arranged in descending order. These values are the square roots of the eigenvalues of both $A^T A$ and AA^T .

4.2 Applications of SVD

SVD has numerous applications:

- It's used in machine learning and data science to perform dimensionality reduction, particularly through a technique known as Principal Component Analysis (PCA).

- In numerical linear algebra, SVD is used to solve linear equations and compute matrix inverses in a more numerically stable way.
- It's used in image compression, where low-rank approximations of an image matrix provide a compressed version of the original image.

4.3 Calculating SVD Manually

You might be inclined to skip this example because the computations are lengthy. Why would anyone do this when they can use a computing language, like Python, to calculate the SVD with essentially one command? We show this so you can understand what goes on "under the hood" when you compute SVD programmatically.

After you read through this example, you'll see how to use Python to compute SVD. Then you'll see an example of using SVD for image compression. Finally, you'll have an exercise to compute the SVD. For this, you'll need to write your own Python script.

Let's find the SVD for matrix A . Recall that we want to find U , Σ , and V^T such that:

$$A = U\Sigma V^T \quad (4.2)$$

$$A = \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix}$$

$$U = AA^T$$

Calculating

$$AA^T$$

will give us a square matrix:

$$A^T = \begin{bmatrix} 3 & -1 \\ 1 & 3 \\ 1 & 1 \end{bmatrix}$$

$$AA^T = \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix} \begin{bmatrix} 3 & -1 \\ 1 & 3 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 11 & 1 \\ 1 & 11 \end{bmatrix}$$

Next we will find the eigenvalues and eigenvectors of A^T . This is a chance to apply what you learned in the previous chapter. We know that:

$$Av = \lambda v \quad (4.3)$$

So:

$$\begin{bmatrix} 11 & 1 \\ 1 & 11 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \lambda \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

Rewrite as a set of equations:

$$11x_1 + x_2 = \lambda x_1$$

$$x_1 + 11x_2 = \lambda x_2$$

Then rearrange:

$$(11-\lambda)x_1 + x_2 = 0$$

$$x_1 + (11-\lambda)x_2 = 0$$

Solve for λ :

$$\begin{bmatrix} (11-\lambda), 1 \\ 1, (11-\lambda) \end{bmatrix} = 0$$

And as equations:

$$(11-\lambda)(11-\lambda) - 1 \cdot 1 = 0$$

$$(\lambda-10)(\lambda-12) = 0$$

These are the eigenvalues.

$$\lambda = 10$$

$$\lambda = 12$$

When substituted into the original equations, you get the eigenvectors. For

$$\lambda = 10$$

:

$$(11-10)x_1 + x_2 = 0$$

$$x_1 = -x_2$$

We'll set

$$x_1$$

to 1 and get this eigenvector:

$$[1, -1]$$

For

$$\lambda = 12$$

:

$$(11-12)x_1 + x_2 = 0$$

$$x_1 = x_2$$

We'll set

$$x_1$$

to 1 and get this eigenvector:

$$[1, 1]$$

The matrix is:

$$\begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

Next you need to apply the Gram-Schmidt process to the column vectors. Then you'll have U , the $m \times m$ matrix whose columns are eigenvectors of AA^T . These are the left singular vectors of A . After you apply Gram-Schmidt, you should end up with:

$$U = \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix}$$

The process for calculating V is the same as the calculation for U , except:

$$V = A^T A$$

$$A^T A = \begin{bmatrix} 3 & -1 \\ 1 & 3 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix} = \begin{bmatrix} 10 & 0 & 2 \\ 0 & 10 & 4 \\ 2 & 4 & 2 \end{bmatrix}$$

After applying the process we applied to solve for U , you get:

$$V = \begin{bmatrix} 1/\sqrt{6} & 2/\sqrt{5} & 1/\sqrt{30} \\ 2/\sqrt{6} & -1/\sqrt{5} & 2/\sqrt{30} \\ 1/\sqrt{6} & 0 & -5/\sqrt{30} \end{bmatrix}$$

However, you want V_T :

$$V_T = \begin{bmatrix} 1/\sqrt{6} & 2/\sqrt{6} & 1/\sqrt{6} \\ 2/\sqrt{5} & -1/\sqrt{5} & 0 \\ 1/\sqrt{30} & 2/\sqrt{30} & -5/\sqrt{30} \end{bmatrix}$$

You have only to calculate Σ , a diagonal matrix that is the same size as A . The diagonal contains the singular values of A , arranged in descending order. They are the square roots of the eigenvalues of both $A^T A$ and AA^T .

Because the non-zero eigenvalues of U are the same as V , let's use the eigenvalues we calculate for U , 10 and 12. Note that Σ will not be of the correct dimension to reconstruct the original matrix unless we add a column. By adding a zero column you'll be able to multiply between U and V :

$$\Sigma = \begin{bmatrix} \sqrt{12} & 0 & 0 \\ 0 & \sqrt{12} & 0 \end{bmatrix}$$

You can check your work by multiplying the decomposed matrices. This should return the original matrix.

$$A = U\Sigma V^T$$

$$= U = \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix} \begin{bmatrix} \sqrt{12} & 0 & 0 \\ 0 & \sqrt{12} & 0 \end{bmatrix} \begin{bmatrix} 1/\sqrt{6} & 2/\sqrt{6} & 1/\sqrt{6} \\ 2/\sqrt{5} & -1/\sqrt{5} & 0 \\ 1/\sqrt{30} & 2/\sqrt{30} & -5/\sqrt{30} \end{bmatrix}$$

$$= \begin{bmatrix} \sqrt{12}/\sqrt{2} & \sqrt{10}/\sqrt{2} & 0 \\ \sqrt{12}/\sqrt{2} & -\sqrt{10}/\sqrt{2} & 0 \end{bmatrix} \begin{bmatrix} 1/\sqrt{6} & 2/\sqrt{6} & 1/\sqrt{6} \\ 2/\sqrt{5} & -1/\sqrt{5} & 0 \\ 1/\sqrt{30} & 2/\sqrt{30} & -5/\sqrt{30} \end{bmatrix}$$

$$= \begin{bmatrix} 3 & 1 & 1 \\ -1 & 3 & 1 \end{bmatrix}$$

4.4 Singular Value Decomposition with Python

Create a file called `vectors_decomposition.py` and enter this code:

```
# Singular-value decomposition
import numpy as np
from numpy import array
from scipy.linalg import svd
from numpy import diag
from numpy import dot
from numpy import zeros

# Define a matrix
A = array([[1, 2], [3, 4], [5, 6]])

print("Matrix (3x2) to be decomposed: ")
print(A)

# CalculateSVD
U, S, VT = svd(A)
print("Matrix (3x3) that represents the left singular values of A:")
print(U)
print("Singular values:")
print(S)
print("Matrix (2x2) that represents the right singular values of A:")
print(VT)

# Check if the decomposition by rebuilding the original matrix
# The singular values must be in an m x n matrix
# Create a zero matrix with the same dimension as A
Sigma = zeros((A.shape[0], A.shape[1]))
# Populate Sigma with n x n diagonal matrix
Sigma[:A.shape[1], :A.shape[1]] = diag(S)
# Reconstruct the original matrix
A_Rebuilt = U.dot(Sigma.dot(VT))
print("Original matrix:")
print(A_Rebuilt)
```

4.5 Sign Ambiguity

You might notice that at times the absolute values in the U and V^T matrices are correct but that the signs vary from what you see as the answer. For example, when you compare a manually calculated SVD with one done in Python the signs might not agree. Both decompositions of A are valid. Both decompositions will satisfy:

$$A = U\Sigma V^T$$

Note that the S diagonal values will always be positive.

The sign ambiguity has implications. For example, when using SVD to compress data, if some of the signs are flipped, the data can have artifacts. At this point in your education, you don't need to concern yourself with it except when you are comparing SVD results for the same matrix.

Exercise 3 Single Value Decomposition


Modify your Python code to calculate SVD for the matrix in the worked out example. Did you arrive at the same answer? Keep in mind that Python will compute square roots and present fractions as decimal. Take a look at the signs for the values in the U and V^T matrices. Are they the same or is this an example of sign ambiguity?

Working Space

Answer on Page 36

4.6 SVD Applied to Image Compression

This image consists of a grid of 20 by 10 pixels, each of which is either black or white.



It's a simple image that has only two types of columns—ideal for data compression. A row is either the first pattern or the second.



We can represent the data as a 20 by 10 matrix whose 200 entries are either 0 for black or 1 for white.

```
[00001000100000001010]
[00001111111111111010]
[00001000100000001010]
[00001000100000001010]
[00001000100000001010]
[00001000100000001010]
[00001000100000001010]
[00001000100000001010]
[00001111111111111010]
[00001000100000001010]
[00001000100000001010]
```

When you perform an SVD on this matrix, there are only two non-zero singular values, 6.79 and 3.72. (You are welcome perform the calculation in Python.) Thus you can represent the matrix as:

$$A = U_1 S_1 V_1 + U_2 S_2 V_2$$

This means there are two u vectors each with 20 entries and two v vectors each with 10 entries, and two singular values. Add those up: $2 \cdot 20 + 2 \cdot 10 + 2 = 62$. This implies that the image can be represented by 62 values instead of 200. If you look back at the image, you can see that there are many dependent columns and very few independent ones.

This is a simple image and a small pixel matrix. But it should give you a sense of how SVD can decompose an image in a way that identifies how much of the image is redundant, and therefor can be compressed.

4.7 Where to Learn More

We Recommend a Singular Value Decomposition. This American Mathematical Society publication focuses the geometry of SVD. What I like about the article is that it shows both

graphically and numerically how SVD can be used for data compression on images and for noise reduction. The data compression example in your workbook is based on this article. <https://www.ams.org/publicoutreach/feature-column/fcarc-svd>

Sign Ambiguity in Singular Value Decomposition (SVD). This is a good article for those who want a deeper understanding of sign ambiguity. <https://www.educative.io/blog/sign-ambiguity-in-singular-value-decomposition>

Singular Value Decomposition Tutorial. This PDF starts by defining points, space, and vectors and works through all the concepts you need to tackle SVD. It is one of the few resources that has a completely worked out example of manually calculating SVD. The example in this chapter is from that tutorial. If you read the entire paper, you'll find it is a good review of the concepts you've studied in previous chapters. <https://rb.gy/j6s0w>

Tackling Difficult Problems: Positive Semidefinite Matrices

With all the computing power available today, you'd think no problem would be too difficult to tackle. But that's not so. There is a category of problems called NP (non-deterministic polynomial time) whose solution is easy to verify but whose computation is difficult to perform because there is no straightforward algorithm and any brute-force method would take too much time. For these problems, all we can do is to develop an efficient algorithm that can find a solution in a reasonable amount of time. While the solution might not be the most optimal, the goal is to find the best solution possible in a short time.

As you learn more about optimization techniques, you'll come across many efficiency algorithms that have been used throughout the years. In the 1990's, the field of optimization changed with the discovery that algorithms based on semidefinite positive matrices can achieve a higher efficiency than seen in the past. Today, there is an entire field of programming—Semidefinite Programming—based on the use of semidefinite positive matrices.

First we'll take a look at what some of the NP problems are. Then we'll describe the intuition behind semidefinite positive matrices. We'll take a look at one NP problem and then introduce the python module to use for solving these problems.

5.1 NP Problems

In the world of mathematics, easy problems are referred to as P, or polynomial time, problems. In simple terms, this means the problem can be solved quickly and it's easy to verify that the solution is correct. Addition, subtraction, division, multiplication, square roots, matrix-vector multiplication, are just a few examples. But NP problems, as stated in the introduction, can't be computed in any reasonable amount of time but solutions are usually easy to verify. The game of Sudoku is one such example. It's easy to verify a correct solution, but writing a generalized algorithm to solve any Sudoku game is an NP problem.

NP problems show up in many other situations, such as:


- Designing robust communication networks
- Scheduling tasks without conflicts

- Managing supply chains
- Detecting patterns in biological networks
- Figuring out subgroups in social networks
- Predicting the structure of proteins

For each of these situations think of large scale problems for which there are many variables. A cloud computing company that provides AI services to thousands of clients must be able to schedule tasks efficiently and in a timely manner as well as manage the power needed for the computers and cooling the data center. Supply chain management is crucial to figuring out how to pick up, transport, and distribute goods to help provide disaster relief. Understanding protein structure is important for designing drugs that can tackle specific conditions. All we can do for each of these situations is to find an optimal solution, but not the definitive solution.

5.2 A Past Approach: Minimizing Errors in Neural Networks


When neural networks were first being developed to recognize things (faces, letters, music, and so on), the goal was to calculate weights between network nodes that would minimize the recognition error. The error space could be visualized as a surface of valleys and hills. The lowest point would have the least error. The idea behind the iterative weight calculations for training the network was to descend down the gradient until reaching the low point. Without getting into the mathematics, you can see by looking at this figure that there are two valleys. Some neural network training resulted in ending at a low point, but not the lowest point. Wouldn't it be great if you could formulate an optimization problem so that you'd be guaranteed to land at the lowest point? That's where semidefinite programming can help.




5.3 Positive Definite and Semidefinite Matrices

Unlike matrices that are defined by their content (such as identity matrix, zero matrix, and diagonal matrix) positive definite and semidefinite matrices are characterized by the result they produce. They have an analog in the scalar world, so let's first look at that. Let's take the scalars a and b and treat them as vectors. ab is then the dot product. If $a > 0$, then ab will take on the sign of b . If $a < 0$, then ab will have the opposite sign of b . If you look at this graphically, you can see that in the first case ab stays on the same side of the origin, but in the second case, ab flips

For $a = [2]$, $b = [4]$



For $a = [3]$, $b = [-4]$. the result of multiplication flips a to the other side of the graph.



The notion of “staying on the same side” is positive definite. The notion of flipping is negative definite. Positive definite means that $x > 0$, so the result is a positive number. Positive semidefinite means that $x \geq 0$, so the result is a non-negative number.

If you can formulate a problem as a positive semidefinite matrix, then you automatically constrain the result to the “same side.” This constraint results in higher algorithmic efficiency.

Let's look at the formal definition:

A matrix is positive definite if, and only if:

$$x^T A x > 0, x, x \neq 0$$

A matrix is positive semidefinite, if, and only if:


$$x^T A x \geq 0$$

Further, a positive semidefinite matrix has an important property. Its eigenvalues are ≥ 0 . The eigenvalues of a positive definite matrix are > 0 .

Take the triplet (a, b, c) and the symmetric matrix:

$$\begin{bmatrix} 1 & a & b \\ a & 1 & c \\ b & c & 1 \end{bmatrix} \geq 0$$

For what values of a, b, c is this matrix positive semidefinite? If you iterate through all possible combinations of a, b, c and then compute the eigenvalues for each matrix, you'll find that some (like $0, 0, 0$) result in a positive semidefinite matrix and some (like $2, 2, 2$) are not positive semidefinite. If you plot the set of triplets that result in a semidefinite matrix, you'll see an ellipsope. It's this shape that guarantees an optimal solution.



5.4 Identifying and Constructing a Positive Semidefinite Matrix

In the last section, you saw that being symmetric does not guarantee a positive semidefinite matrix. You also saw that a matrix of positive values does not guarantee a positive semidefinite matrix. The only way to check for positive semidefinite is to calculate the eigenvalues, which you learned in a previous workbook.

A surefire way to construct a positive semidefinite matrix is:

$$AA^T$$

Exercise 4 Figuring out if a matrix is positive semidefinite

Is this matrix positive definite? Show your work.

$$\begin{bmatrix} 2 & 2 \\ 2 & 0 \end{bmatrix}$$

Working Space

Answer on Page 36

Exercise 5 Creating a positive semidefinite matrix

Using any 3 by 3 matrix, create a positive semidefinite matrix. Then show it is positive semidefinite by calculating its eigenvalues. You can either compute this by hand or using python. In either case, show your work.

Working Space

Answer on Page 36

5.5 The Max Cut Problem

A famous NP problem is Max Cut. Given a graph of interconnected nodes, cut the graph to create two sets of nodes, such that the cut goes through as many edges as possible. (You can't cut an edge more than once.) Max Cut is important for binary classification in machine learning, circuit design, statistical physics, and more. There is no algorithm that will provide an exact solution. If you could find one, you'd be eligible to win a huge prize from the Clay Mathematics Institute. Instead, you'll see how to approximate a solution to this problem using a positive semidefinite matrix and a technique developed by mathematicians Michel Goemans and David Williamson.


You won't see all the complete details here as this section is meant to be a quick introduction to how you can apply positive semidefinite matrices.

Take this simple graph of five nodes and six edges. Each node in the graph will take on

one of two values (1 or -1) to show which set they fall into after a cut is made. For any two connected nodes, $x_i * x_j = 1$ if $x_i = x_j$ and -1 otherwise.

If you randomly assign 1 and -1 to the nodes, the chance of making the max cut is 0.5. By using semidefinite programming you can achieve an algorithmic efficiency of 0.87.

The Goemans-Williamson technique can be used for any optimization problem where the variables take on the values of 1 and -1 .



As a list of edges the graph is:

$$\text{edges} = [(0, 1), (0, 2), (1, 3), (1, 4), (2, 3), (3, 4)]$$

The optimization problem can be formulated as:

$$\text{Max} \sum_{\text{edges}(i,j)} \frac{1 - x_i x_j}{2}$$

for

$$x_i \in \{-1, 1\}$$

BUT instead of allowing x_i to be scalar, Goemans-Williamson defines x_i as unit vectors.

$$x_i \in \mathbb{R}^n, \|x_i\| = 1$$

and that makes the optimization equation:

$$\text{Max} \sum_{\text{edges}(i,j)} \frac{1 - x_i^T x_j}{2}$$

It's this "relaxation" that gets us to a semidefinite matrix, because we can now rewrite the problem as a positive semidefinite matrix:

$$X = [x_i^T x_j]_{i,j}$$

Python has a module for solving optimization problems. Using this, you'll get an optimum

matrix, but to get the unit vectors, you'll need to take the square root of the matrix.


$$X = [x_1 \dots x_{1n}] [x_1 \dots x_{1n}]^T$$

Next we need to go from unit vectors to scalars, a process called rounding.

$$x^i \in \mathbb{R}^n \rightarrow x^i \in \{-1, 1\}$$

Goemans-Williamson leveraged the fact that the end point of a unit vector is on a sphere. They generated a random plane to bisect the sphere. A vector on one side of the plane is assigned the value of 1 and a vector on the other side a value of -1 .

You can then assign the scalar values to the nodes and make the cut accordingly. This particular cut will be 5, as shown.



5.6 The Max Cut Problem Solved in Python

```
# MaxCut Problem
import numpy as np
import scipy
from scipy.linalg import sqrtm
# cvxpy is a python module for solving optimization problems
import cvxpy as cp

# define the edges of the graph
edges = [(0,1),
         (0,2),
         (1,3),
         (1,4),
         (2,3),
         (3,4)]

# Declare the matrix X to be positive semidefinite
X = cp.Variable((5,5),symmetric=True)
```

```
constraints = [X >> 0]

# Set diagonals to 1 to get unit vectors
constraints += [
    X[i,i] == 1 for i in range(5)
]

# Set the objective function
objective = sum(0.5*(1 - X[i,j]) for (i,j) in edges)

# Set up the problem to maximize using the objective function and
# keeping within the set constraints
prob = cp.Problem(cp.Maximize(objective), constraints)

# Returns a positive semidefinite matrix
print(prob.solve())

# To get the vectors take square root of the matrix
x = sqrtm(X.value)

# Generate a random hyperplane
u = np.random.randn(5) # normal to random hyperplane

# Pick values according to which side of the hyperplane the vectors are on
x = np.sign(x @ u)
```

Answers to Exercises

Answer to Exercise 1 (on page 6)

Compute dot product of **a** and **b**:

$$1 * -4 + 3 * 6 = -4 + 18 = 14$$

Compute the dot product of **b** and **b**

$$16 + 36 = 52$$

$$14/52 * (-4, 6) = (-1.076, 1.61)$$

Answer to Exercise 2 (on page 10)

The first vector of the orthogonal subspace is:

$$v_1 = x_1 = (1, 1, 1)$$

The second vector of the subspace is a projection of x_2 onto v_1 .

$$v_2 = x_2 - \frac{x_2 v_1}{v_1 v_1} v_1$$

Substitute the values:

$$v_2 = (0, 1, 1) - \frac{(0, 1, 1)(1, 1, 1)}{(1, 1, 1)(1, 1, -1)}(1, 1, 1)$$

$$v_2 = (0, 1, 1) - (2/3)(1, 1, 1)$$

$$v_2 = (-2/3, 1/3, 1/3)$$

Normalize:

$$v_1 = v_1 / \sqrt{|v_1|}$$

$$v_1 = (1, 1, 1) / \sqrt{|v_1|}$$

$$v_1 = (0.577, 0.577, 0.577)$$

$$v_2 = v_2 / \sqrt{|v_2|}$$

$$\begin{aligned}v_2 &= (0, 1, 1)\sqrt{|v_2|} \\v_2 &= (-0.816, 0.408, 0.408)\end{aligned}$$

Answer to Exercise 3 (on page 24)

$$\begin{aligned}U &= \begin{bmatrix} -0.70710678 & -0.70710678 \\ -0.70710678 & 0.70710678 \end{bmatrix} \\ \text{Singularvalues} &= [3.464101623.16227766] \\ V^T &= \begin{bmatrix} -0.408 & -0.816 & -0.408 \\ -0.894 & 0.447 & 0.0 \\ -0.183 & -0.365 & 0.9129 \end{bmatrix}\end{aligned}$$

Answer to Exercise 4 (on page 31)

Yes. It's eigenvalues are

$$2, 2$$

Answer to Exercise 5 (on page 31)

The answer depends on the 3 by 3 matrix you chose.



INDEX

eigenvalue, [13](#)

eigenvector, [13](#)

Gram-Schmidt process, [7](#)

projection, [3](#)

singular value decomposition, [19](#)

svd, [19](#)